IS459 BDA Assignment 2

Question 1

How large are the communities (connected components)?

The graph used here to analyze the connected components is formed by connecting the nodes of each author (source author) with the node of another author (destination author) that they interacted with on the HardwareZone PC Gaming Forum. As a form of data pre-processing, I've also removed the edges where the source author = destination author, which was created when an author comments on his/her own post *(lines 54-58 of assignment_2.py)*.

By applying the connectedComponents() function on the graph, we notice that out of 4661 distinct authors on HardwareZone, 4536 of them are interconnected with each other, and there are many other smaller connected components with 3 or less authors in it, as seen in this DataFrame below. *(lines 75-80 of assignment_2.py)*

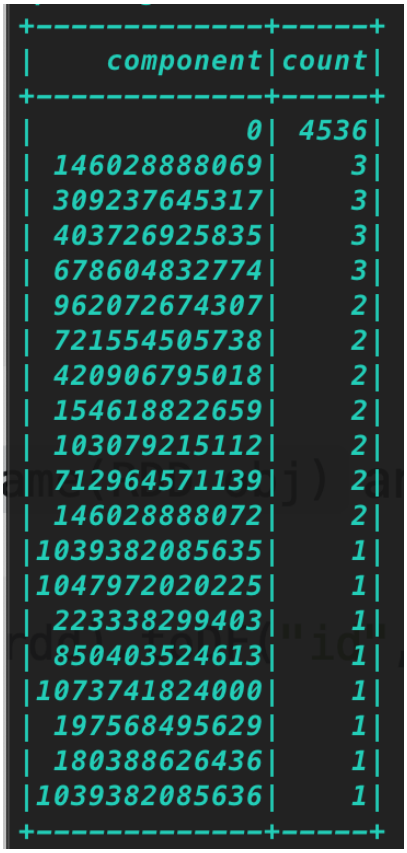

component	count
0	4536
146028888069	3
309237645317	3
403726925835	3
678604832774	3
962072674307	2
721554505738	2
420906795018	2
154618822659	2
103079215112	2
712964571139	2
146028888072	2
1039382085635	1
1047972020225	1
223338299403	1
850403524613	1
1073741824000	1
197568495629	1
180388626436	1
1039382085636	1

As such, I would conclude that in general, the 4661 authors on HardwareZone are rather **closely connected with one another** as a majority (4536 authors) are connected to one another in one single connected component. There are a total of **111** connected components in the community, where **110** of these connected components contain **less than or equal to 3 authors**.

What are the key words of the community (frequent words)?

In order to get the frequent keywords, I've retrieved a list of stopwords from https://countwordsfree.com/stopwords and downloaded it as a text file, and uploaded it to HDFS so that I'm able to read it on assignment_2.py. By tokenizing every word in content and removing it if it's a stopword, I've gathered the following list of common keywords, sorted in descending order.

Attached below is a list of the top 50 keywords, sorted in descending order.

(lines 82-96 of assignment_2.py)

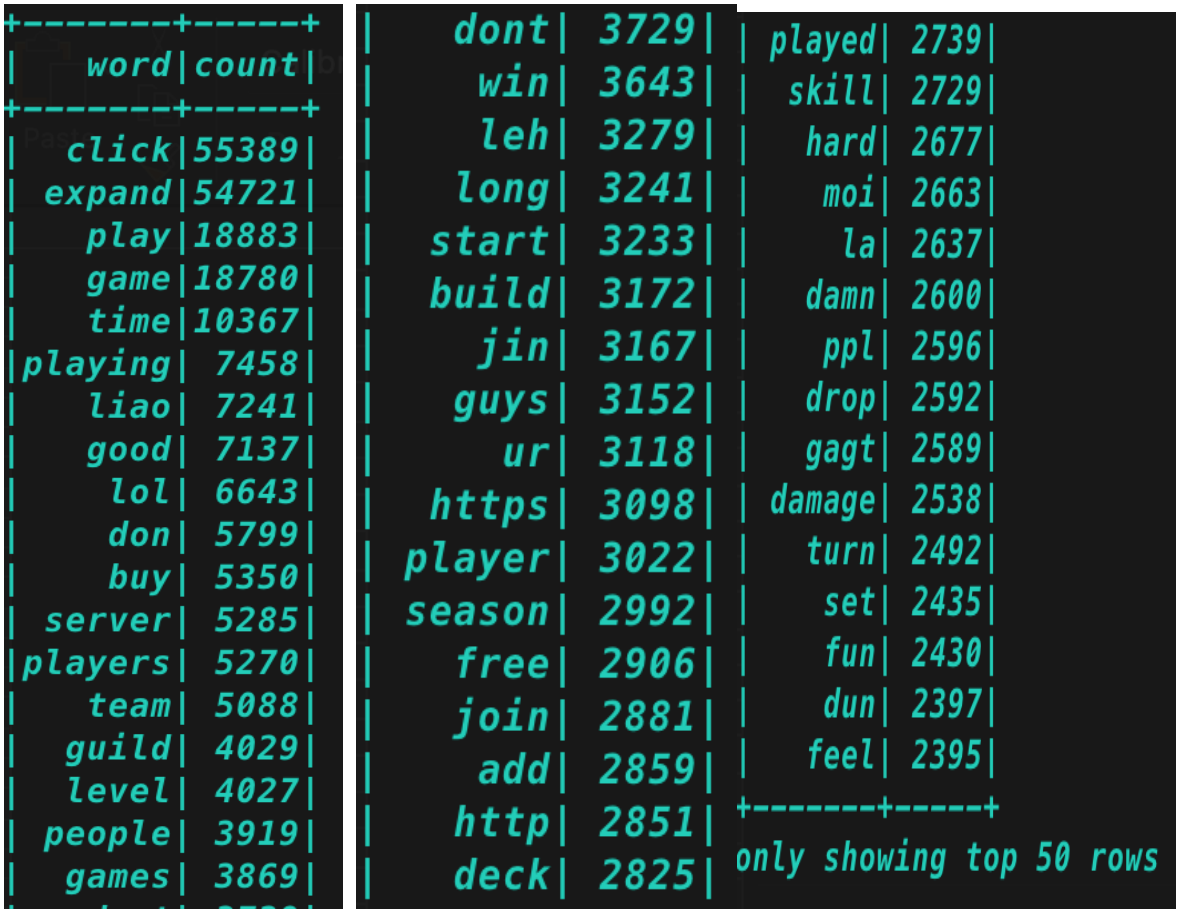

word	count
click	55389
expand	54721
play	18883
game	18780
time	10367
playing	7458
liao	7241
good	7137
lol	6643
don	5799
buy	5350
server	5285
players	5270
team	5088
guild	4029
level	4027
people	3919
games	3869
dont	3729
win	3643
leh	3279
long	3241
start	3233
build	3172
jin	3167
guys	3152
ur	3118
https	3098
player	3022
season	2992
free	2906
join	2881
add	2859
http	2851
deck	2825
played	2739
skill	2729
hard	2677
moi	2663
la	2637
damn	2600
ppl	2596
drop	2592
gagt	2589
damage	2538
turn	2492
set	2435
fun	2430
dun	2397
feel	2395

only showing top 50 rows

Question 2

How cohesive are the communities (Average # of triangles over every user in a community)?

(lines 101-106 of assignment_2.py) By applying the triangleCount() function on the graph, we're able to get the total number of triangles formed for each author. We can then sum up the number of triangles, which is **264211899**.

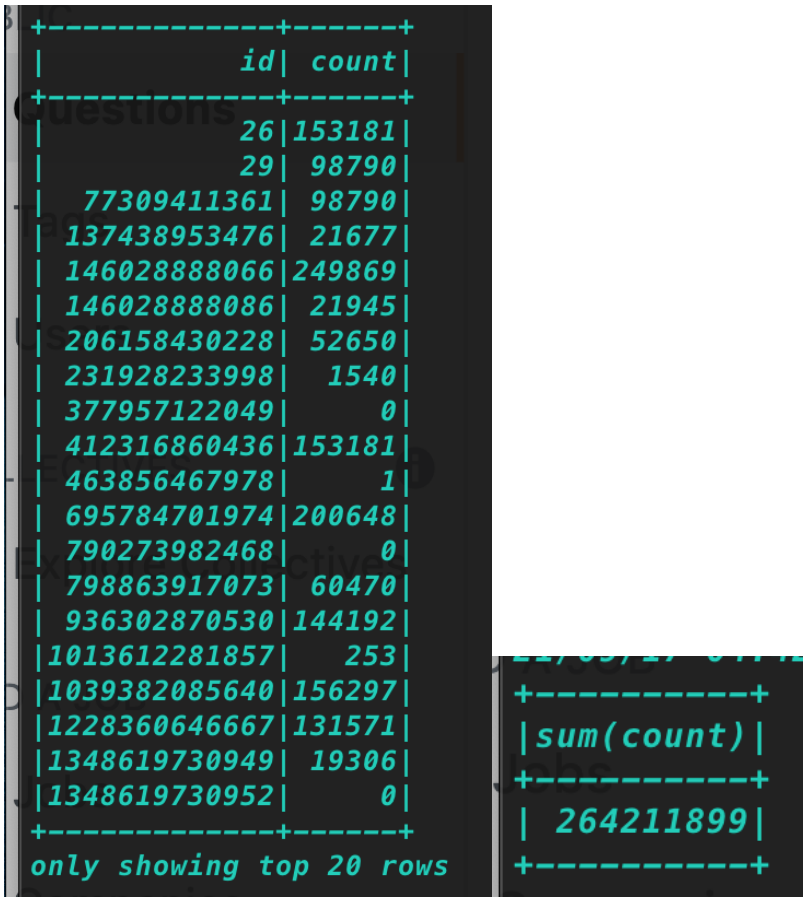

```
+------------+------+
|          id| count|
+------------+------+
|          26|153181|
|          29| 98790|
|  77309411361| 98790|
| 137438953476| 21677|
| 146028888066|249869|
| 146028888086| 21945|
| 206158430228| 52650|
| 231928233998|  1540|
| 377957122049|     0|
| 412316860436|153181|
| 463856467978|     1|
| 695784701974|200648|
| 790273982468|     0|
| 798863917073| 60470|
| 936302870530|144192|
|1013612281857|   253|
|1039382085640|156297|
|1228360646667|131571|
|1348619730949| 19306|
|1348619730952|     0|
+------------+------+
only showing top 20 rows
```

```
+----------+
|sum(count)|
+----------+
| 264211899|
+----------+
```

In order to determine the cohesiveness of the community, we can then find out the average number of triangles across all the authors on the HardwareZone PC Gaming Forum, which is an average of **56685** triangles. Hence, we can conclude that the community in the HardwareZone PC Gaming Forum is indeed cohesive.

```
>>> totalTriangles
264211899
>>> avgTriangles = totalTriangles / author_df.count()
>>> avgTriangles
56685.6680969749
```

Question 3

Is there any strange community? (Open) – E.g., younger generation or older generation?

(lines 111-120 of assignment_2.py) For this question, I've crawled additional data of the Join Date of every author on HardwareZone. The Join Date can also be a good gauge of the "age" of the author on the HardwareZone PC Gaming Forum, and for this question, we'll be making the following assumptions:

- An author is considered younger generation if he/she joined after 2011 (Year >= 2011)
- An author is considered older generation if he/she joined before 2011 (Year < 2011)

First, we can gather the data and create a dataframe as seen below (left image), with a data frame of the author and the year they've joined HardwareZone.

author	year
pebblesontheway	2021
BalaKype	2021
CircuitBreakerKia	2021
Lummyz	2021
semc88	2021
bigdaddy1234	2021
Leftyaof	2021
hhkjss	2021
bendoggo	2021
Mortgage Advisor	2021
goldtoes68	2021
REIT-FI	2021
Gss2021	2021
malabuyaola	2021
Chavvo	2021
mocha_latte	2021
Ajrail	2021
maychua	2021
LubbyLub	2021
MeSoFancy	2020

only showing top 20 rows

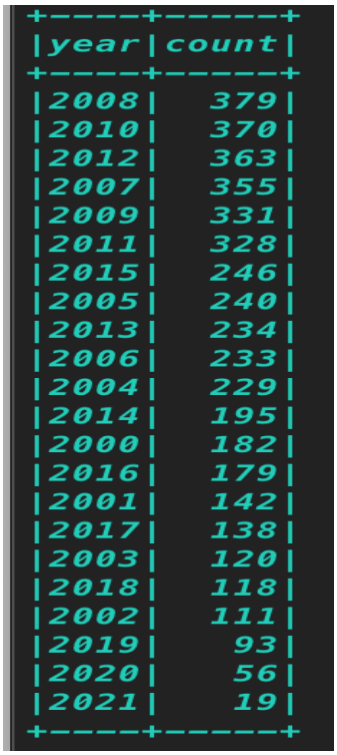

year	count
2008	379
2010	370
2012	363
2007	355
2009	331
2011	328
2015	246
2005	240
2013	234
2006	233
2004	229
2014	195
2000	182
2016	179
2001	142
2017	138
2003	120
2018	118
2002	111
2019	93
2020	56
2021	19

Next, we can then find out the number of authors who joined every year and create a dataframe for that as well (right image).

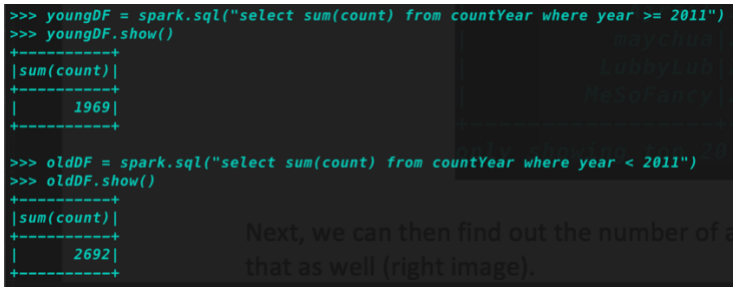

```
>>> youngDF = spark.sql("select sum(count) from countYear where year >= 2011")
>>> youngDF.show()
+----------+
|sum(count)|
+----------+
|      1969|
+----------+

>>> oldDF = spark.sql("select sum(count) from countYear where year < 2011")
>>> oldDF.show()
+----------+
|sum(count)|
+----------+
|      2692|
+----------+
```

From the image above, we can tell that there's actually more authors under the older generation category **(2692 authors**), as compared to the younger generation category (**1969 authors**). Hence, I'll conclude that there's **more authors from the older generation**, based on the assumption mentioned above.